

 PULSE Rec
Dashboard

- View
- ☐

 A/B/C Analysis
- ☒

 Advanced Eval
- ☐

 Single run
- ☐

 Compare runs
- ☐

 All runs timeline
- ☐

 Raw logs

8 result file(s) loaded



Advanced Evaluation


Three data-driven tests on existing log data — no new API calls. Run `python scripts/eval_advanced.py` to generate or refresh results.

Showing: `advanced_eval_20260430T233142Z` — Places run `20260430T051138Z` , Yelp run `20260430T213155Z`

Test 1 — Cross-source head-to-head


Same 15 queries fired at Google Places and Yelp Fusion. Client re-ranker applied independently to each pool, then to a merged pool.

Places avg NDCG@5	Yelp avg NDCG@5	Merged avg NDCG@5	Head-to-head
0.9670	0.9290	0.9387	P 6W · Y 3...
↓ -0.0380			



Top-5 slot attribution (merged pool)

How many of the merged top-5 slots come from each source.




	ID	Query	Places NDCG@5	Yelp NDCG@5	Merged NDCG@5	Win
0	q01	best tapas near me	1.000000	0.830400	0.830400	place
1	q02	Italian dinner tonight	1.000000	1.000000	1.000000	tie
2	q03	yoga class tomorrow morning	1.000000	1.000000	1.000000	tie
3	q04	craft beer bar	1.000000	0.915200	1.000000	place
4	q05	brunch this weekend	0.830400	0.927000	0.861400	yelp
5	q06	contemporary art museum	1.000000	0.915200	0.915200	place
6	q07	hidden gem restaurant	0.853900	0.893000	0.868800	yelp
7	q08	coffee shop to work from	0.893000	1.000000	0.893000	yelp
8	q09	rooftop bar	0.927000	0.849600	0.849600	place
9	q10	japanese ramen	1.000000	1.000000	1.000000	tie

Test 2 — Re-ranker signal ablation

Zero out one signal at a time on Places data. Negative delta = that signal helps. Cold-start signals (saved, prefTags, netAffinity) are always 0 — ablating them has no effect.

Baseline NDCG@5 (full scorer)

0.9670



Red bars (negative delta) = signal **helps** ranking — removing it hurts. Green bar (positive delta) = signal **hurts** ranking — removing it helps. 🚩 `isOpen` (++) is counterproductive: venues marked closed get penalized even if they are high quality, so NDCG@5 improves when you stop penalizing them.

Active signals

Cold-start signals (always 0 — no personalization)

	Signal	Weight	NDCG@5 (ablated)		Signal	Weight	Note
1	reviews	0.120000	0.9413	3	saved	0.14	always 0 at cold-start — abl
6	queryAffinity	0.140000	0.9419	5	preferredTags	0.08	always 0 at cold-start — abl
7	distance	0.070000	0.9516	8	networkAffinity	0.06	always 0 at cold-start — abl
0	rating	0.240000	0.956600				
4	preferredCat	0.110000	0.967000				
2	isOpen	0.140000	0.979300				

Test 3 — Intra-list Diversity (ILD@5)

ILD = fraction of top-5 pairs with different categories. 0.0 = all 5 results same category, 1.0 = all different. Diversity is good for multi-intent queries (q11, q14) but precision queries (q02 Italian dinner) should score near 0.


Places avg ILD@5

0.1333

Yelp avg ILD@5

0.2067

↑ +0.0734



	ID	Query	Places ILD	Yelp ILD	Places cat dist
0	q01	best tapas near me	0.000000	0.400000	{'food': 5}
1	q02	Italian dinner tonight	0.000000	0.000000	{'food': 5}
2	q03	yoga class tomorrow morning	0.000000	0.000000	{'wellness': 5}
3	q04	craft beer bar	0.000000	0.400000	{'nightlife': 5}
4	q05	brunch this weekend	0.400000	0.000000	{'nightlife': 1, 'food': 4}
5	q06	contemporary art museum	0.000000	0.000000	{'culture': 5}
6	q07	hidden gem restaurant	0.400000	0.400000	{'food': 4, 'nightlife': 1}
7	q08	coffee shop to work from	0.400000	0.000000	{'food': 4, 'wellness': 1}
8	q09	rooftop bar	0.400000	0.000000	{'nightlife': 4, 'food': 1}
9	q10	japanese ramen	0.000000	0.000000	{'food': 5}

Advanced eval run: [advanced_eval_20260430T233142Z](#) — Places run [20260430T051138Z](#) · Yelp run [20260430T213155Z](#)